

# Reinforcement Learning Basics

**Shusen Wang**

Stevens Institute of Technology

<http://wangshusen.github.io/>

**A little bit probability theory...**

# Random Variable

- **Random variable**: unknown; its values depend on outcomes of random events.

*Random  
Variable*

*Possible  
Values*

*Random  
Events*

*Probabilities*

$$X = \begin{cases} 0 \\ 1 \end{cases}$$



$$\mathbb{P}(X = 0) = 0.5$$

$$\mathbb{P}(X = 1) = 0.5$$

# Random Variable

- **Random variable**: unknown; its values depend on outcomes of random events.
- Uppercase letter  $X$  for a random variable.
- Lowercase letter  $x$  for an observed value.
- For example, I flipped a coin 4 times and observed:
  - $x_1 = 1$ ,
  - $x_2 = 1$ ,
  - $x_3 = 0$ ,
  - $x_4 = 1$ .

# Probability Density Function (PDF)

- PDF provides a relative likelihood that the value of the random variable would equal that sample.


# Probability Density Function (PDF)

- PDF provides a relative likelihood that the value of the random variable would equal that sample.

## Example: Gaussian distribution

- It is a continuous distribution.
- PDF:

$$p(x) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right).$$



# Probability Mass Function (PMF)

- PMF is a function that gives the probability that a discrete random variable is exactly equal to some value.


# Probability Mass Function (PMF)

- PMF is a function that gives the probability that a discrete random variable is exactly equal to some value.

## Example

- Discrete random variable:  $X \in \{1, 3, 7\}$ .
- PDF:

$$\begin{aligned}p(1) &= 0.2, \\p(3) &= 0.5, \\p(7) &= 0.3.\end{aligned}$$





# Probability Mass Function (PMF)

- PMF is a function that gives the probability that a discrete random variable is exactly equal to some value.


## Example

- Discrete random variable:  $X \in \{1, 3, 7\}$ .
- PDF:

$$p(1) = 0.2,$$

$$p(3) = 0.5,$$

$$p(7) = 0.3.$$



# Properties of PDF/PMF

- Random variable  $X$  is in the domain  $\mathcal{X}$ .

- For continuous distributions,

$$\int_{\mathcal{X}} p(x) dx = 1.$$

- For discrete distributions,

$$\sum_{x \in \mathcal{X}} p(x) = 1.$$

# Expectation

- Random variable  $X$  is in the domain  $\mathcal{X}$ .
- For continuous distributions, the expectation of  $f(X)$  is:


$$\underline{\mathbb{E} [f(X)]} = \underline{\int_{\mathcal{X}} p(x) \cdot f(x) dx}.$$

- For discrete distributions, the expectation of  $f(X)$  is:

$$\underline{\mathbb{E} [f(X)]} = \underline{\sum_{x \in \mathcal{X}} p(x) \cdot f(x)}.$$

# Random Sampling

- There are 10 balls in the bin: 2 are red, 5 are green, and 3 are blue.



# Random Sampling

- Sample red ball w.p. 0.2, green ball w.p. 0.5, and blue ball w.p. 0.3.

# Random Sampling

- Sample red ball w.p. 0.2, green ball w.p. 0.5, and blue ball w.p. 0.3.

```
from numpy.random import choice
```

```
samples = choice(['R', 'G', 'B'], size=100, p=[0.2, 0.5, 0.3])
```

```
print(samples)
```


```
[ 'R'  'G'  'R'  'R'  'R'  'R'  'B'  'B'  'B'  'G'  'G'  'B'  'G'  'B'  'B'  'G'  'B'  'G'
  'B'  'B'  'G'  'B'  'G'  'B'  'B'  'G'  'B'  'B'  'G'  'B'  'G'  'G'  'G'  'G'  'G'  'B'
  'B'  'B'  'B'  'B'  'B'  'G'  'G'  'B'  'R'  'R'  'B'  'R'  'B'  'G'  'R'  'G'  'R'  'G'
  'R'  'R'  'B'  'G'  'G'  'G'  'B'  'R'  'G'  'B'  'G'  'R'  'G'  'G'  'G'  'B'  'B'  'R'
  'G'  'G'  'B'  'B'  'R'  'B'  'B'  'B'  'R'  'B'  'G'  'B'  'R'  'B'  'R'  'G'  'B'  'R'
  'B'  'B'  'G'  'G'  'G'  'R'  'R'  'B'  'R'  'G' ]
```

# Terminologies

# Terminology: state and action

state  $s$  (this frame)

Action  $a \in \{\text{left}, \text{right}, \text{up}\}$






# Terminology: state and action

state  $s$  (this frame)



Action  $a \in \{\text{left}, \text{right}, \text{up}\}$




# Terminology: policy

policy  $\pi$

- Policy function  $\pi: (s, a) \mapsto [0,1]$ :

$$\pi(a | s) = \mathbb{P}(A = a | S = s).$$




# Terminology: policy

policy  $\pi$

- Policy function  $\pi: (s, a) \mapsto [0,1]$ :

$$\pi(a | s) = \mathbb{P}(A = a | S = s).$$






# Terminology: policy

policy  $\pi$


- Policy function  $\pi: (s, a) \mapsto [0,1]$ :

$$\pi(a \mid s) = \mathbb{P}(A = a \mid S = s).$$



# Terminology: policy


policy  $\pi$



- $\pi(a | s)$  is the probability of taking action  $A = a$  given state  $s$ , e.g.,
  - $\pi(\text{left} | s) = 0.2$ ,
  - $\pi(\text{right} | s) = 0.1$ ,
  - $\pi(\text{up} | s) = 0.7$ .

# Terminology: policy

policy  $\pi$




- $\pi(a | s)$  is the probability of taking action  $A = a$  given state  $s$ , e.g.,
  - $\pi(\text{left} | s) = 0.2$ ,
  - $\pi(\text{right} | s) = 0.1$ ,
  - $\pi(\text{up} | s) = 0.7$ .
- Upon observing state  $S = s$ , the agent's action  $A$  can be random.



# Terminology: policy

Random or deterministic policy?



# Terminology: reward

reward  $R$

- Collect a coin:  $R = +1$





# Terminology: reward

reward  $R$



- Collect a coin:  $R = +1$
- Win the game:  $R = +10000$

# Terminology: reward

## reward $R$



- Collect a coin:  $R = +1$
- Win the game:  $R = +10000$
- Touch a Goomba:  $R = -10000$  (game over).

# Terminology: reward

## reward $R$



- Collect a coin:  $R = +1$
- Win the game:  $R = +10000$
- Touch a Goomba:  $R = -10000$  (game over).
- Nothing happens:  $R = 0$



# Terminology: state transition




state transition




# Terminology: state transition

## state transition



- E.g., “up” action leads to a new state.

# Terminology: state transition



## state transition



- E.g., “up” action leads to a new state.
- State transition can be random.
- Randomness is from the environment.



# Terminology: state transition




## state transition



- E.g., “up” action leads to a new state.
- State transition can be random.
- Randomness is from the environment.

# Terminology: state transition



## state transition




- E.g., “up” action leads to a new state.
- State transition can be random.
- Randomness is from the environment.
- $\underline{p(s'|s, a)} = \mathbb{P}(S' = s' | S = s, A = a)$ .




# Two Sources of Randomness

# Randomness in Actions




# Randomness in Actions



Given state  $s$ , the action can be random, e.g., .

- $\pi(\text{"left"} | s) = 0.2,$
- $\pi(\text{"right"} | s) = 0.1,$
- $\pi(\text{"up"} | s) = 0.7.$


# Randomness in States



- State transition can be random.
- The environment generates the new state  $s'$  by

$$\underline{s' \sim p(\cdot | s, a) .}$$

# Randomness in States



- State transition can be random.
- The environment generates the new state  $s'$  by

$$s' \sim p(\cdot | s, a) .$$

# Two Sources of Randomness

- The randomness in **action** is from the policy function:


$$A \sim \pi(\cdot \mid s) .$$

- The randomness in **state** is from the state-transition function:

$$S' \sim p(\cdot \mid s, a) .$$

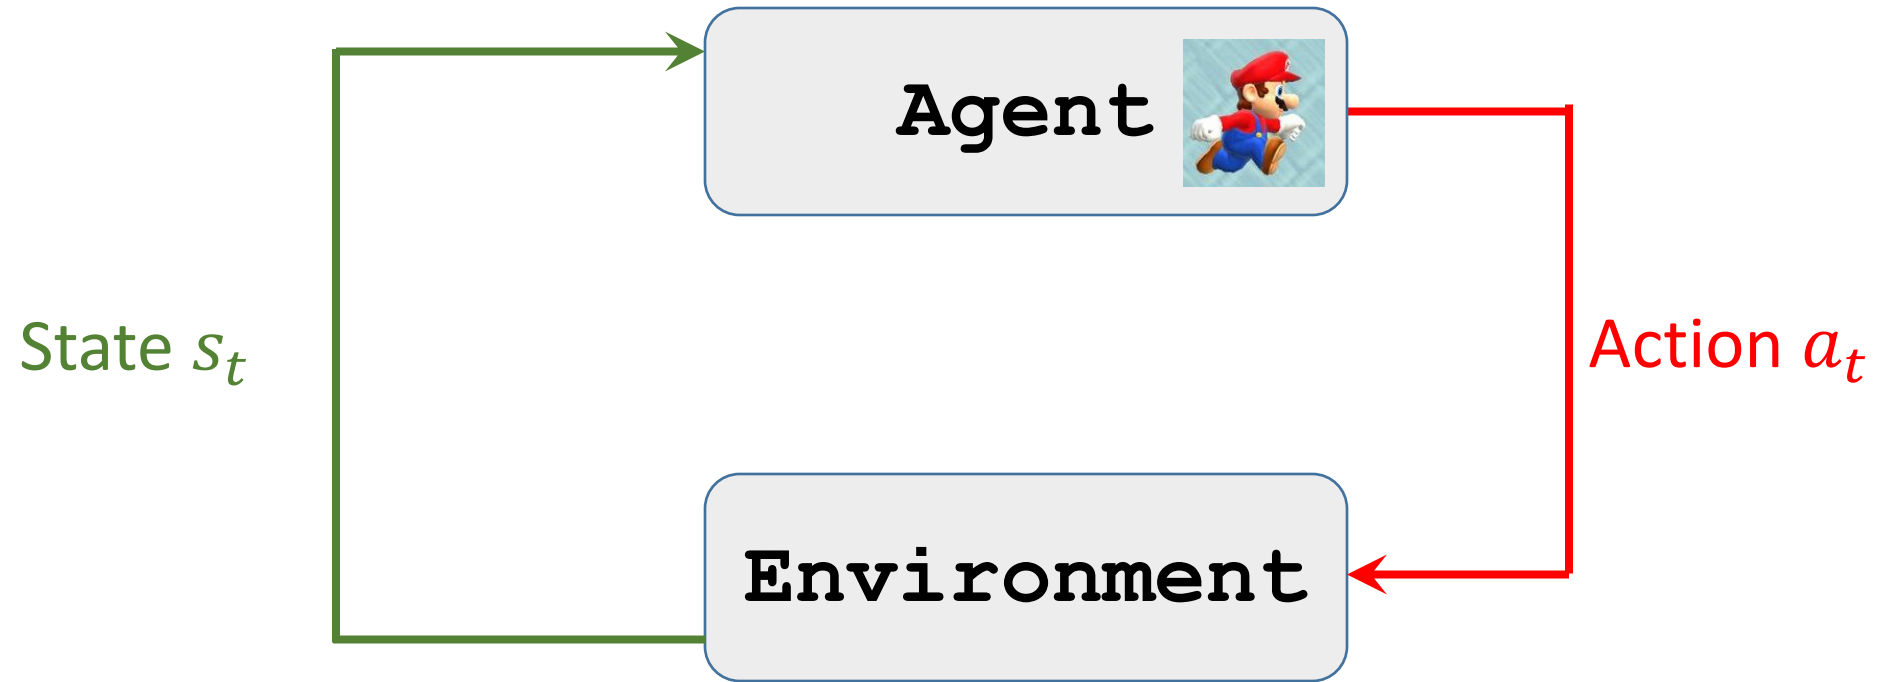
# **Agent-Environment Interaction**

# Agent-Environment Interaction

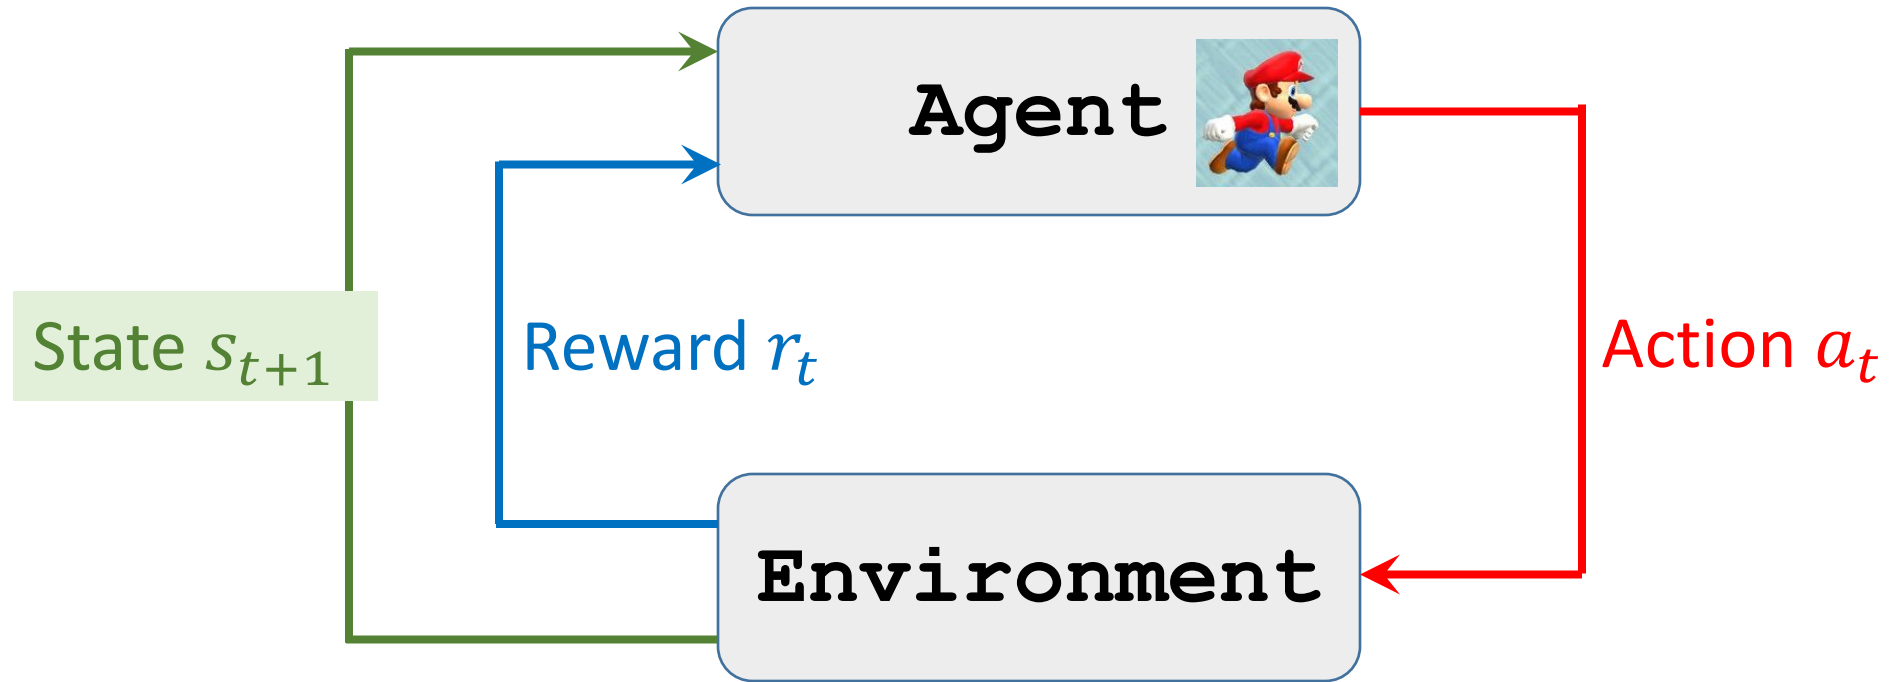




# Agent-Environment Interaction



# Agent-Environment Interaction




# Play game using AI

- Observe state  $s_t$ , select action  $a_t \sim \pi(\cdot \mid s_t)$ , and execute  $a_t$ .
- The environment gives new state  $s_{t+1}$  and reward  $r_t$ .


# Play game using AI

- Observe state  $s_t$ , select action  $a_t \sim \pi(\cdot \mid s_t)$ , and execute  $a_t$ .
- The environment gives new state  $s_{t+1}$  and reward  $r_t$ .



# Play game using AI

- Observe state  $s_t$ , select action  $a_t \sim \pi(\cdot | s_t)$ , and execute  $a_t$ .
- The environment gives new state  $s_{t+1}$  and reward  $r_t$ .




# Play game using AI

- (state, action, reward) trajectory:

$s_1, a_1, r_1, s_2, a_2, r_2, \dots, s_n, a_n, r_n$

- One episode is from the the beginning to the end (Mario wins or dies).



# Play game using AI

- (state, action, reward) trajectory:

$$s_1, a_1, r_1, \quad s_2, a_2, r_2, \quad \cdots, \quad s_n, a_n, r_n.$$

- One episode is from the the beginning to the end (Mario wins or dies).
- What is a good policy?
- A good policy leads to big cumulative reward:  $\sum_{t=1}^n \gamma^{t-1} \cdot r_t$ .
- Use the rewards to guide the learning of policy.

# Rewards and Returns



# Return

**Definition:** Return (aka cumulative future reward).

- $\underline{U_t} = \underline{R_t + R_{t+1} + R_{t+2} + R_{t+3} + \dots}$

# Return

**Definition:** Return (aka cumulative future reward).

- $U_t = R_t + R_{t+1} + R_{t+2} + R_{t+3} + \dots$

**Question:** At time  $t$ , are  $R_t$  and  $R_{t+1}$  equally important?

- Which of the followings do you prefer?
  - I give you \$100 right now.
  - I will give you \$100 one year later.

# Return

**Definition:** Return (aka cumulative future reward).

- $U_t = R_t + R_{t+1} + R_{t+2} + R_{t+3} + \dots$

**Question:** At time  $t$ , are  $R_t$  and  $R_{t+1}$  equally important?

- Which of the followings do you prefer?
  - I give you \$80 right now.
  - I will give you \$100 one year later.

# Return

**Definition:** Return (aka cumulative future reward).

- $U_t = R_t + R_{t+1} + R_{t+2} + R_{t+3} + \dots$

**Question:** At time  $t$ , are  $R_t$  and  $R_{t+1}$  equally important?

- Which of the followings do you prefer?
  - I give you \$80 right now.
  - I will give you \$100 one year later.
- Future reward is less valuable than present reward.
- $R_{t+1}$  should be given less weight than  $R_t$ .

# Discounted Returns

**Definition:** Return (aka cumulative future reward).

- $U_t = R_t + R_{t+1} + R_{t+2} + R_{t+3} + \dots$

**Definition:** Discounted return (aka cumulative discounted future reward).

- $\gamma$ : discount factor (tuning hyper-parameter).

- $U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \gamma^3 R_{t+3} + \dots$

# Discounted Returns

**Definition:** Discounted return (at time  $t$ ).

- $U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \dots + \gamma^{n-t} R_n.$

# Randomness in Returns

**Definition:** Discounted return (at time  $t$ ).

- $U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \dots + \gamma^{n-t} R_n.$

At the end of the game, we observe  $u_t$ .

- We observe all the rewards,  $r_t, r_{t+1}, r_{t+2}, \dots, r_n.$
- We thereby know the discounted return:

$$u_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots + \gamma^{n-t} r_n.$$

# Randomness in Returns

**Definition:** Discounted return (at time  $t$ ).

- $U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \dots + \gamma^{n-t} R_n.$

At time  $t$ , the rewards,  $R_t, \dots, R_n$ , are **random**, so the return  $U_t$  is **random**.



# Randomness in Returns

**Definition:** Discounted return (at time  $t$ ).

- $U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \dots + \gamma^{n-t} R_n.$

At time  $t$ , the rewards,  $R_t, \dots, R_n$ , are **random**, so the return  $U_t$  is **random**.

- Reward  $R_i$  depends on  $S_i$  and  $A_i$ .
- States can be random:  $S_i \sim p(\cdot \mid s_{i-1}, a_{i-1}).$
- Actions can be random:  $A_i \sim \pi(\cdot \mid s_i).$
- If either  $S_i$  or  $A_i$  is random, then  $R_i$  is random.

# Randomness in Returns


**Definition:** Discounted return (at time  $t$ ).

- $U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \dots + \gamma^{n-t} R_n.$

At time  $t$ , the rewards,  $R_t, \dots, R_n$ , are **random**, so the return  $U_t$  is **random**.

- Reward  $R_i$  depends on  $S_i$  and  $A_i$ .
- $U_t$  depends on  $R_t, R_{t+1}, \dots, R_n$ .
- $\Rightarrow U_t$  depends on  $S_t, A_t, S_{t+1}, A_{t+1}, \dots, S_n, A_n$ .

# Randomness in Returns



# Randomness in Returns



$R_t$

$R_{t+1}$

$R_{t+2}$

...

$R_n$

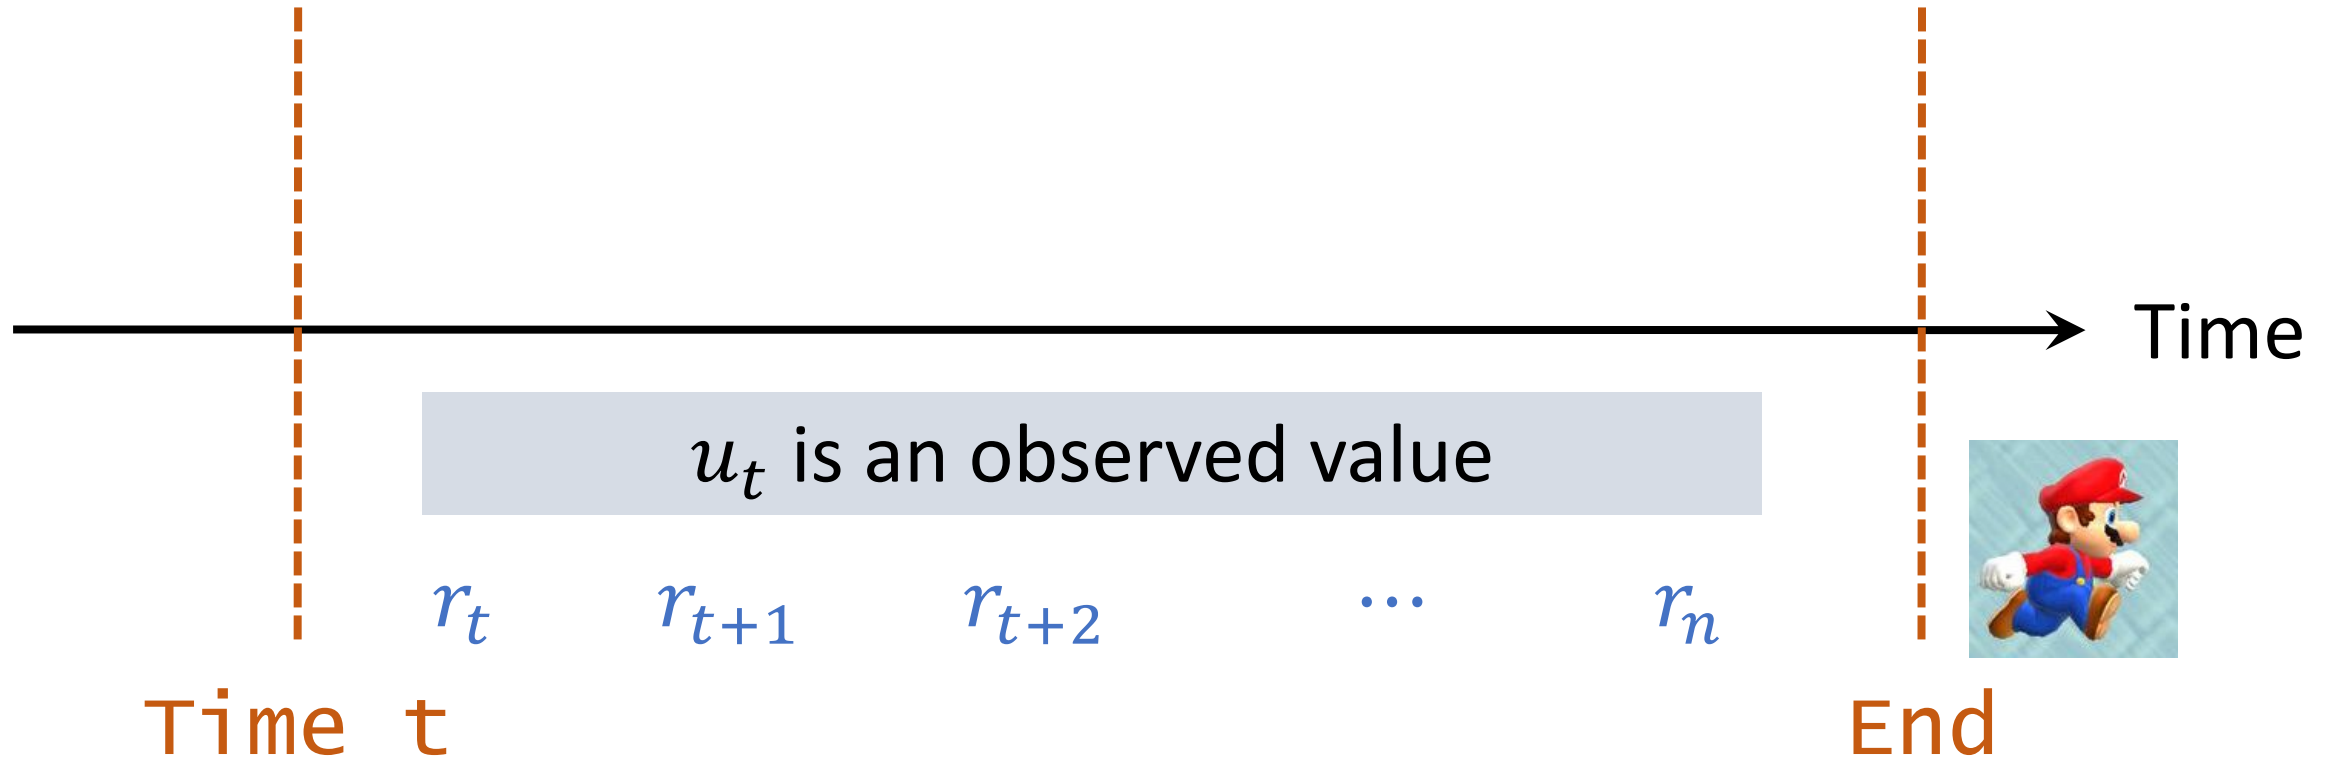
$U_t$  is a random variable

Time

Time t

End

# Randomness in Returns



# **Value Functions**

# Action-Value Function $Q_{\pi}(s, a)$

**Definition:** Discounted return.

- $U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \dots + \gamma^{n-t} R_n.$

# Action-Value Function $Q_{\pi}(s, a)$

**Definition:** Discounted return.

- $U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \dots + \gamma^{n-t} R_n.$

**Definition:** Action-value function.

- $Q_{\pi}(s_t, a_t) = \mathbb{E} [ U_t \mid S_t = s_t, A_t = a_t ].$



# Action-Value Function $Q_{\pi}(s, a)$

**Definition:** Discounted return.

- $U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \dots + \gamma^{n-t} R_n.$

**Definition:** Action-value function.

- $Q_{\pi}(s_t, a_t) = \mathbb{E} [ U_t \mid S_t = s_t, A_t = a_t ].$



$U_t$  depends on states  $s_t, s_{t+1}, \dots, s_n$  and actions  $a_t, a_{t+1}, \dots, a_n$ .

# Action-Value Function $Q_\pi(s, a)$

**Definition:** Discounted return.

- $U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \dots + \gamma^{n-t} R_n.$

**Definition:** Action-value function.

- $Q_\pi(s_t, a_t) = \mathbb{E} [ U_t \mid S_t = s_t, A_t = a_t ].$

Regard  $s_t$  and  $a_t$  as observed values.

Regard  $\underline{s_{t+1}, \dots, s_n}$  and  $\underline{A_{t+1}, \dots, A_n}$  as random variables.

# Action-Value Function $Q_{\pi}(s, a)$

**Definition:** Discounted return.

- $U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \dots + \gamma^{n-t} R_n.$

**Definition:** Action-value function.

- $Q_{\pi}(s_t, a_t) = \mathbb{E}[U_t \mid S_t = s_t, A_t = a_t].$

Regard  $s_{t+1}, \dots, s_n$  and  $A_{t+1}, \dots, A_n$  as random variables.

# Action-Value Function $Q_{\pi}(s, a)$

**Definition:** Discounted return.

$$\bullet U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \dots + \gamma^{n-t} R_n.$$

**Definition:** Action-value function.

$$\bullet Q_{\pi}(s_t, a_t) = \mathbb{E}[U_t \mid S_t = s_t, A_t = a_t].$$

$$\begin{aligned} \bullet S_{t+1} &\sim p(\cdot \mid s_t, a_t), \\ &\vdots \\ \bullet S_n &\sim p(\cdot \mid s_{n-1}, a_{n-1}). \end{aligned}$$

$$\begin{aligned} \bullet A_{t+1} &\sim \pi(\cdot \mid s_{t+1}), \\ &\vdots \\ \bullet A_n &\sim \pi(\cdot \mid s_n). \end{aligned}$$

# Action-Value Function $Q_{\pi}(s, a)$

**Definition:** Discounted return.

- $U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \dots + \gamma^{n-t} R_n.$

**Definition:** Action-value function.

- $Q_{\pi}(s_t, a_t) = \mathbb{E} [ U_t \mid S_t = s_t, A_t = a_t ].$



# Action-Value Function $Q_\pi(s, a)$

**Definition:** Discounted return.

- $U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \dots + \gamma^{n-t} R_n.$

**Definition:** Action-value function.

- $Q_\pi(\underline{s_t}, a_t) = \mathbb{E} [ U_t \mid S_t = s_t, A_t = a_t ].$

# Action-Value Function $Q_{\pi}(s, a)$

**Definition:** Discounted return.

- $U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \dots + \gamma^{n-t} R_n.$

**Definition:** Action-value function.

- $Q_{\pi}(s_t, a_t) = \mathbb{E} [ U_t \mid S_t = s_t, A_t = a_t ].$

- $Q_{\pi}(s_t, a_t)$  depends on  $s_t, a_t, \pi$ , and  $p$ .

- $Q_{\pi}(s_t, a_t)$  is dependent of  $\underline{s_{t+1}, \dots, s_n}$  and  $\underline{A_{t+1}, \dots, A_n}.$

# State-Value Function $V_{\pi}(s)$

**Definition:** Discounted return.

- $U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \dots + \gamma^{n-t} R_n.$

**Definition:** Action-value function.

- $Q_{\pi}(s_t, a_t) = \mathbb{E} [ U_t \mid S_t = s_t, A_t = a_t ].$

**Definition:** State-value function.

- $\underline{V_{\pi}(s_t)} = \mathbb{E}_{\underline{A}} [\underline{Q_{\pi}(s_t, A)}]$



# State-Value Function $V_{\pi}(s)$

**Definition:** Discounted return.

- $U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \dots + \gamma^{n-t} R_n.$

**Definition:** Action-value function.

- $Q_{\pi}(s_t, a_t) = \mathbb{E} [ U_t \mid S_t = s_t, A_t = a_t ].$

**Definition:** State-value function.

- $V_{\pi}(s_t) = \mathbb{E}_A [Q_{\pi}(s_t, A)]$

Taken w.r.t. the action  $A \sim \pi(\cdot \mid s_t).$

# State-Value Function $V_\pi(s)$

**Definition:** Discounted return.

- $U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \dots + \gamma^{n-t} R_n.$

**Definition:** Action-value function.

- $Q_\pi(s_t, a_t) = \mathbb{E} [ U_t \mid S_t = s_t, A_t = a_t ].$

**Definition:** State-value function.

- $V_\pi(s_t) = \mathbb{E}_A [ Q_\pi(s_t, A) ] = \sum_a \pi(a|s_t) \cdot Q_\pi(s_t, a). \quad (\text{Actions are discrete.})$

Taken w.r.t. the action  $A \sim \pi(\cdot | s_t).$

# State-Value Function $V_{\pi}(s)$

**Definition:** Discounted return.

- $U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \dots + \gamma^{n-t} R_n.$

**Definition:** Action-value function.

- $Q_{\pi}(s_t, a_t) = \mathbb{E} [ U_t \mid S_t = s_t, A_t = a_t ].$

**Definition:** State-value function.

- $V_{\pi}(s_t) = \mathbb{E}_A [Q_{\pi}(s_t, A)] = \sum_a \pi(a|s_t) \cdot Q_{\pi}(s_t, a). \quad (\text{Actions are discrete.})$
- $V_{\pi}(s_t) = \mathbb{E}_A [Q_{\pi}(s_t, A)] = \int \pi(a|s_t) \cdot Q_{\pi}(s_t, a) da. \quad (\text{Actions are continuous.})$

# Understanding the Value Functions

- **Action**-value function:  $\underline{Q_\pi(s, a)} = \mathbb{E} [\underline{U_t | S_t = s, A_t = a}]$ .
- Given policy  $\pi$ ,  $Q_\pi(s, a)$  evaluates how good it is for an agent to pick action  $a$  while being in state  $s$ .

# Understanding the Value Functions

- **Action**-value function:  $Q_{\pi}(s, a) = \mathbb{E} [U_t | S_t = s, A_t = a]$ .
- Given policy  $\pi$ ,  $Q_{\pi}(s, a)$  evaluates how good it is for an agent to pick action  $a$  while being in state  $s$ .
- **State**-value function:  $V_{\pi}(s) = \mathbb{E}_A [Q_{\pi}(s, A)]$
- For fixed policy  $\pi$ ,  $V_{\pi}(s)$  evaluates how good the situation is in state  $s$ .
- $\mathbb{E}_S [V_{\pi}(S)]$  evaluates how good the policy  $\pi$  is.

# Evaluating Reinforcement Learning


# OpenAI Gym

- Gym is a toolkit for developing and comparing reinforcement learning algorithms.
- <https://gym.openai.com/>

# OpenAI Gym

- Gym is a toolkit for developing and comparing reinforcement learning algorithms.
- <https://gym.openai.com/>

## Classical control problems



Cart Pole




Pendulum



# OpenAI Gym

- Gym is a toolkit for developing and comparing reinforcement learning algorithms.
- <https://gym.openai.com/>

## Atari Games



Pong



Space Invader




Breakout


# OpenAI Gym

- Gym is a toolkit for developing and comparing reinforcement learning algorithms.
- <https://gym.openai.com/>


## MuJoCo (Continuous control tasks.)



Ant




Humanoid



Half Cheetah


# OpenAI Gym



# Play CartPole Game

```
➡ import gym  
➡ env = gym.make( 'CartPole-v0' )
```

- Get the environment of CartPole from Gym.
- “env” provides states and reward.



# Play CartPole Game

➡ `state = env.reset()`

➡ `for t in range(100):` ➡ A window pops up rendering CartPole.

➡ `env.render()`  
`print(state)`

A random **action**.

➡ `action = env.action_space.sample()`

➡ `state, reward, done, info = env.step(action)`

➡ `if done:` "done=1" means finished (win or lose the game)

`print('Finished')`  
`break`

`env.close()`

# Summary

# Summary

## Terminologies

- Agent 
- Environment
- State  $s$
- Action  $a$
- Reward  $r$

# Summary


## Terminologies

- Agent 
- Environment
- State  $s$
- Action  $a$
- Reward  $r$
- Policy  $\pi(a|s)$
- State transition  $p(s'|s, a)$



# Summary

## Terminologies

- Agent 
- Environment
- State  $s$
- Action  $a$
- Reward  $r$
- Policy  $\pi(a|s)$
- State transition  $p(s'|s, a)$


## Return and Value

- Return:

$$U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \dots$$

# Summary

## Terminologies

- Agent 
- Environment
- State  $s$
- Action  $a$
- Reward  $r$
- Policy  $\pi(a|s)$
- State transition  $p(s'|s, a)$

## Return and Value

- Return:

$$U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \dots$$


- Action-value function:

$$Q_\pi(s_t, a_t) = \mathbb{E}[U_t | s_t, a_t].$$

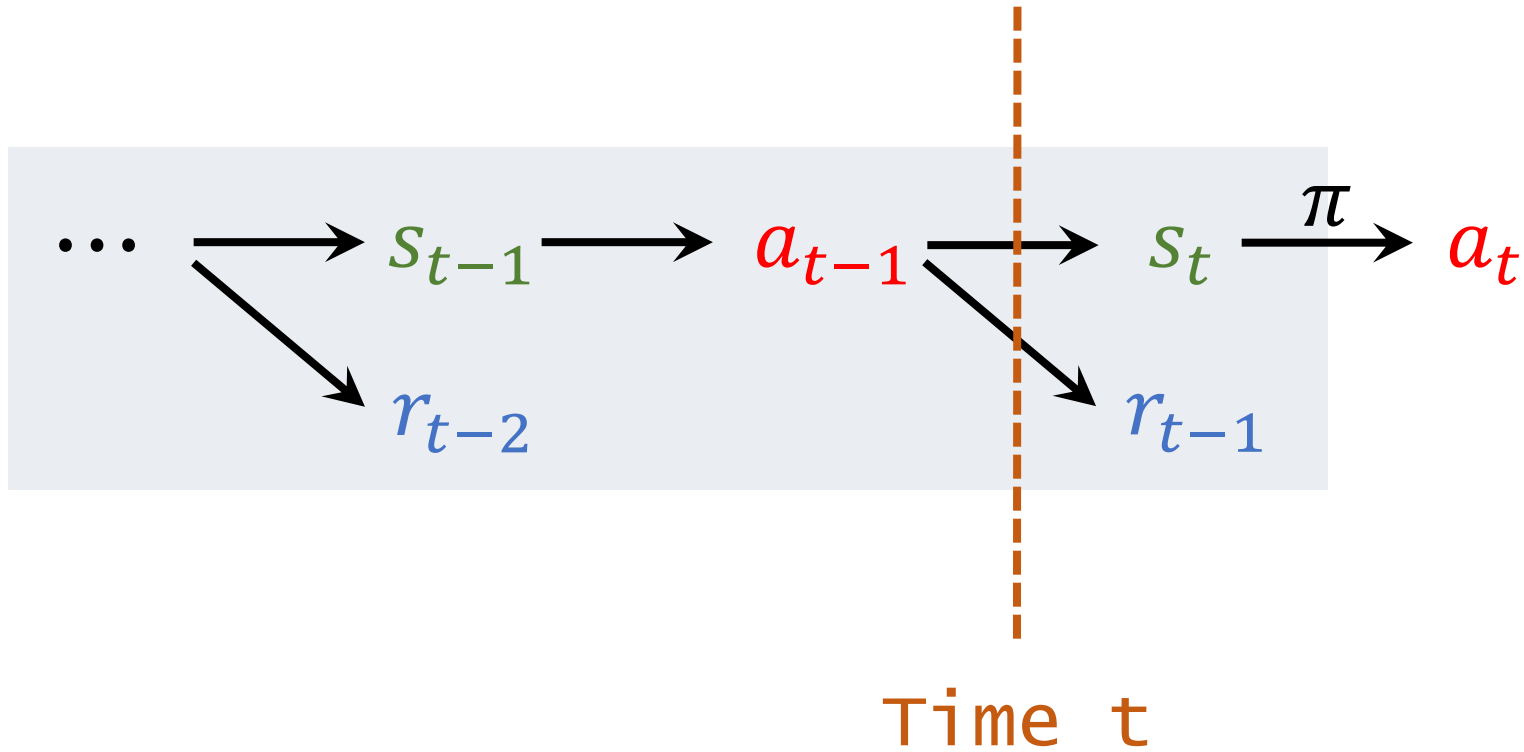
- State-value function:

$$V_\pi(s_t) = \mathbb{E}_{A \sim \pi}[Q_\pi(s_t, A)].$$

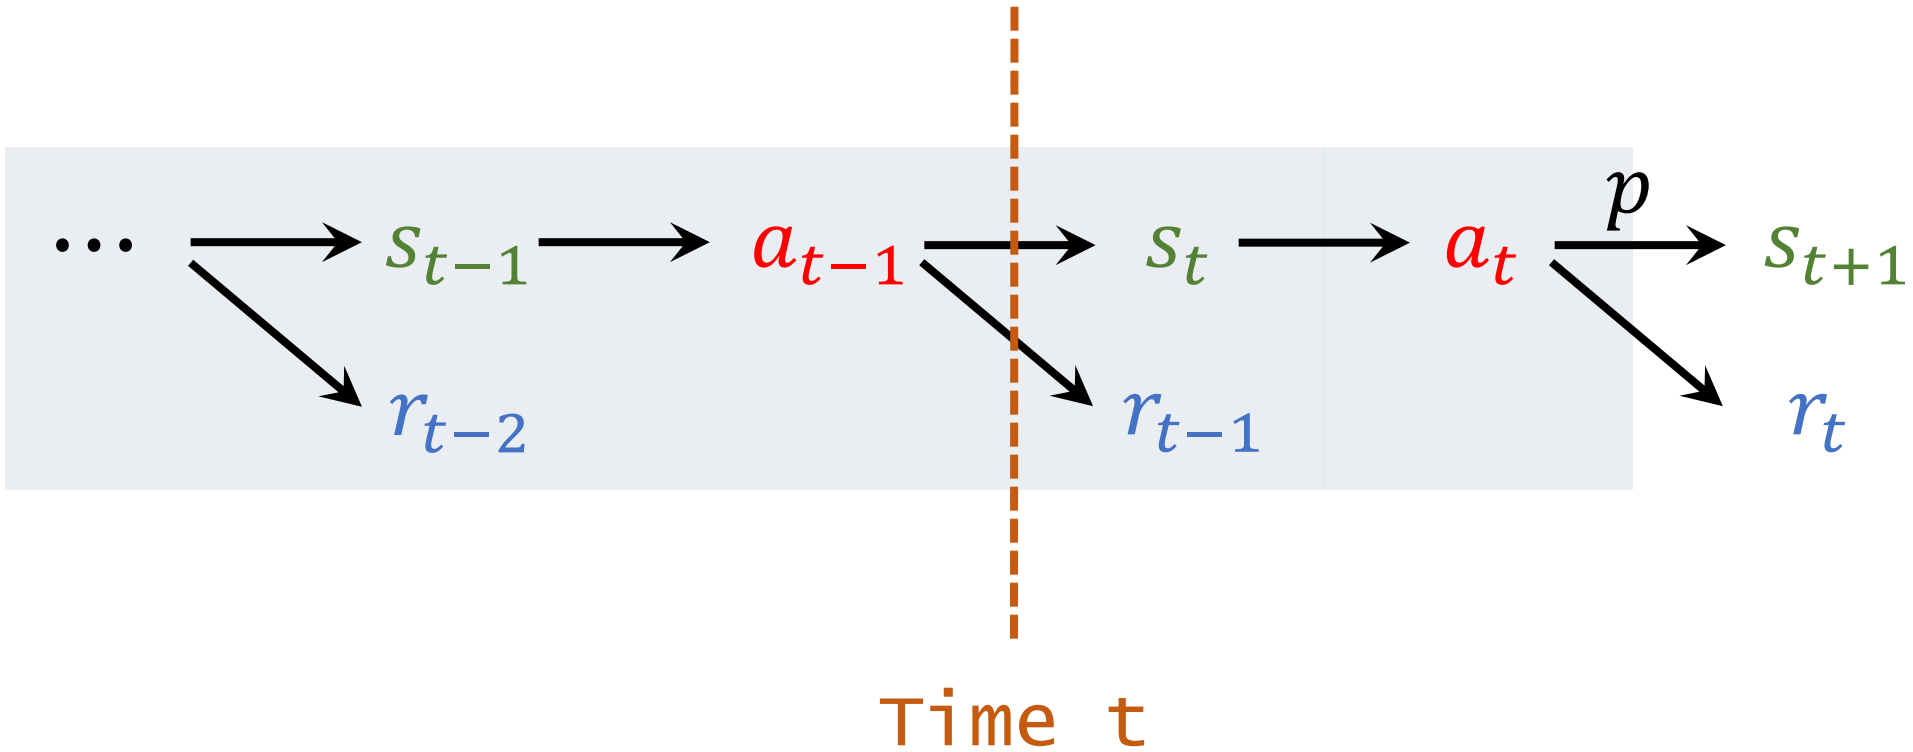
# Play game using AI



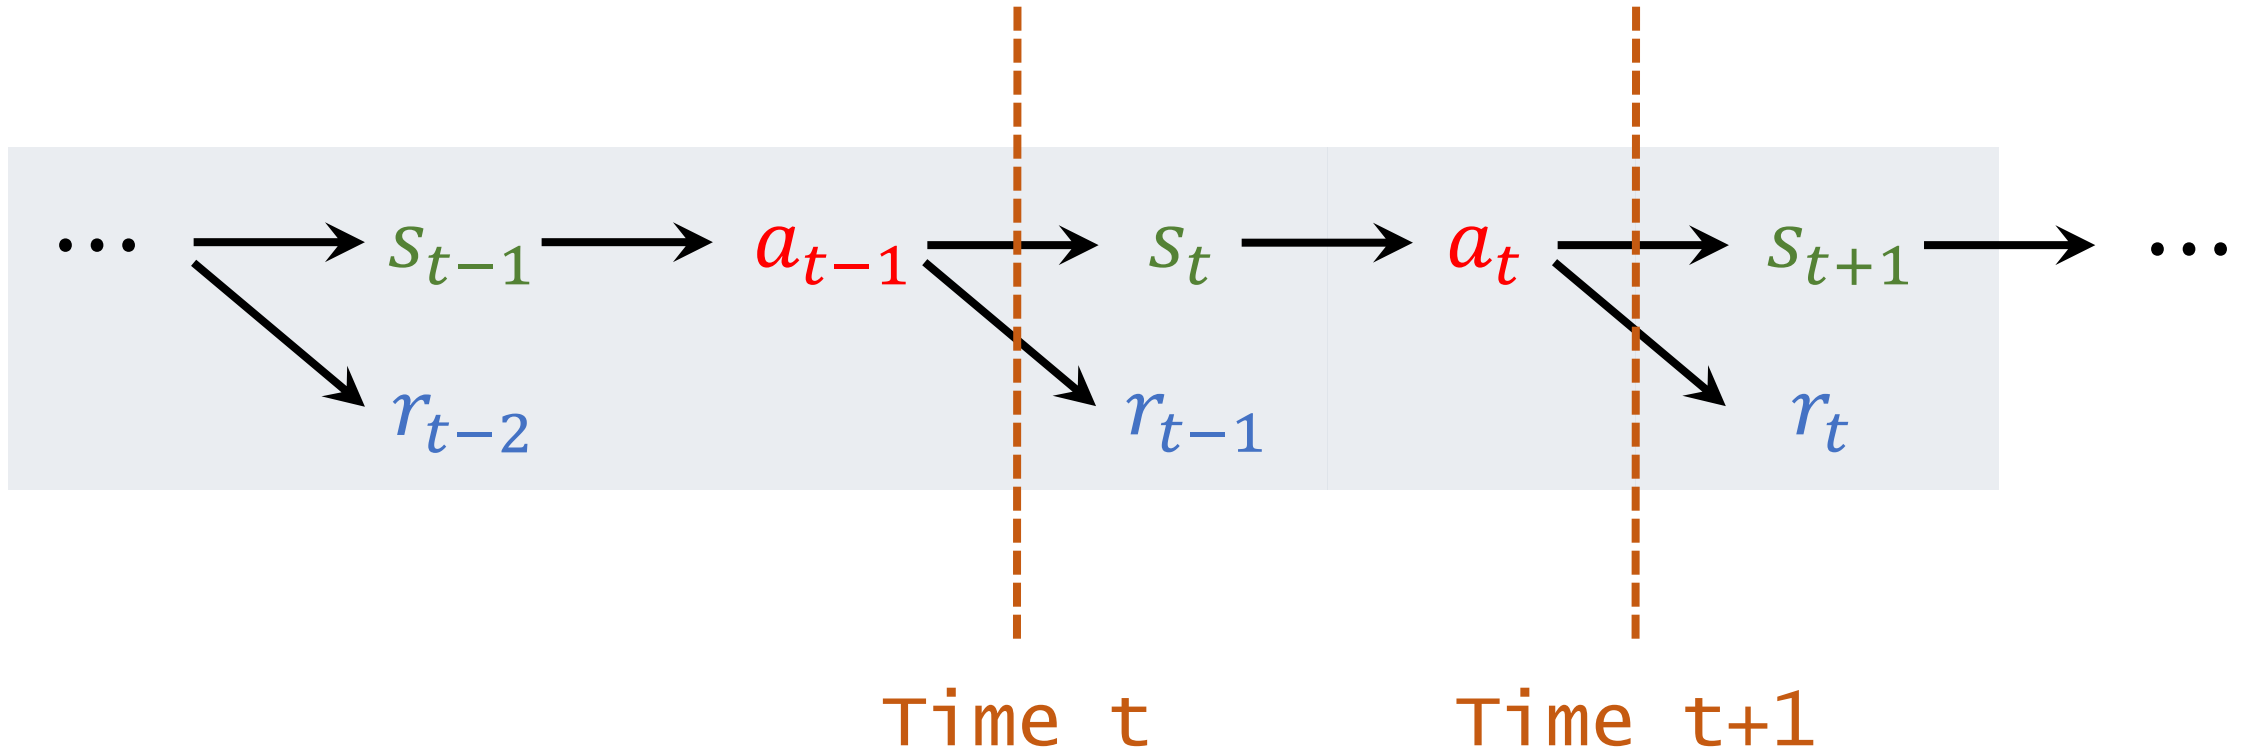
# Play game using AI



# Play game using AI



# Play game using AI



# Thank You!

<http://wangshusen.github.io/>