

# **Policy-Based Reinforcement Learning**

**Shusen Wang**

# **Policy Function Approximation**

# Policy Function $\pi(a|s)$

- Policy function  $\pi(a|s)$  is a probability density function (PDF).
- It takes state  $s$  as input.
- It output the probabilities for all the actions, e.g.,

$$\pi(\text{left}|s) = 0.2,$$

$$\pi(\text{right}|s) = 0.1,$$

$$\pi(\text{up}|s) = 0.7.$$

- Randomly sample action  $a$  random drawn from the distribution.

# Can we directly learn a policy function $\pi(\textcolor{red}{a}|\textcolor{green}{s})$ ?

- If there are only a few states and actions, then yes, we can.
- Draw a table (matrix) and learn the entries.

	Action $a_1$	Action $a_2$	Action $a_3$	Action $a_4$	...
State $s_1$					
State $s_2$					
State $s_3$					
⋮					

# Can we directly learn a policy function $\pi(\textcolor{red}{a}|\textcolor{green}{s})$ ?

- If there are only a few states and actions, then yes, we can.
- Draw a table (matrix) and learn the entries.
- What if there are too many (or infinite) states or actions?

	Action $a_1$	Action $a_2$	Action $a_3$	Action $a_4$	...
State $s_1$					
State $s_2$					
State $s_3$					
⋮					

# Policy Network $\pi(a|s; \theta)$

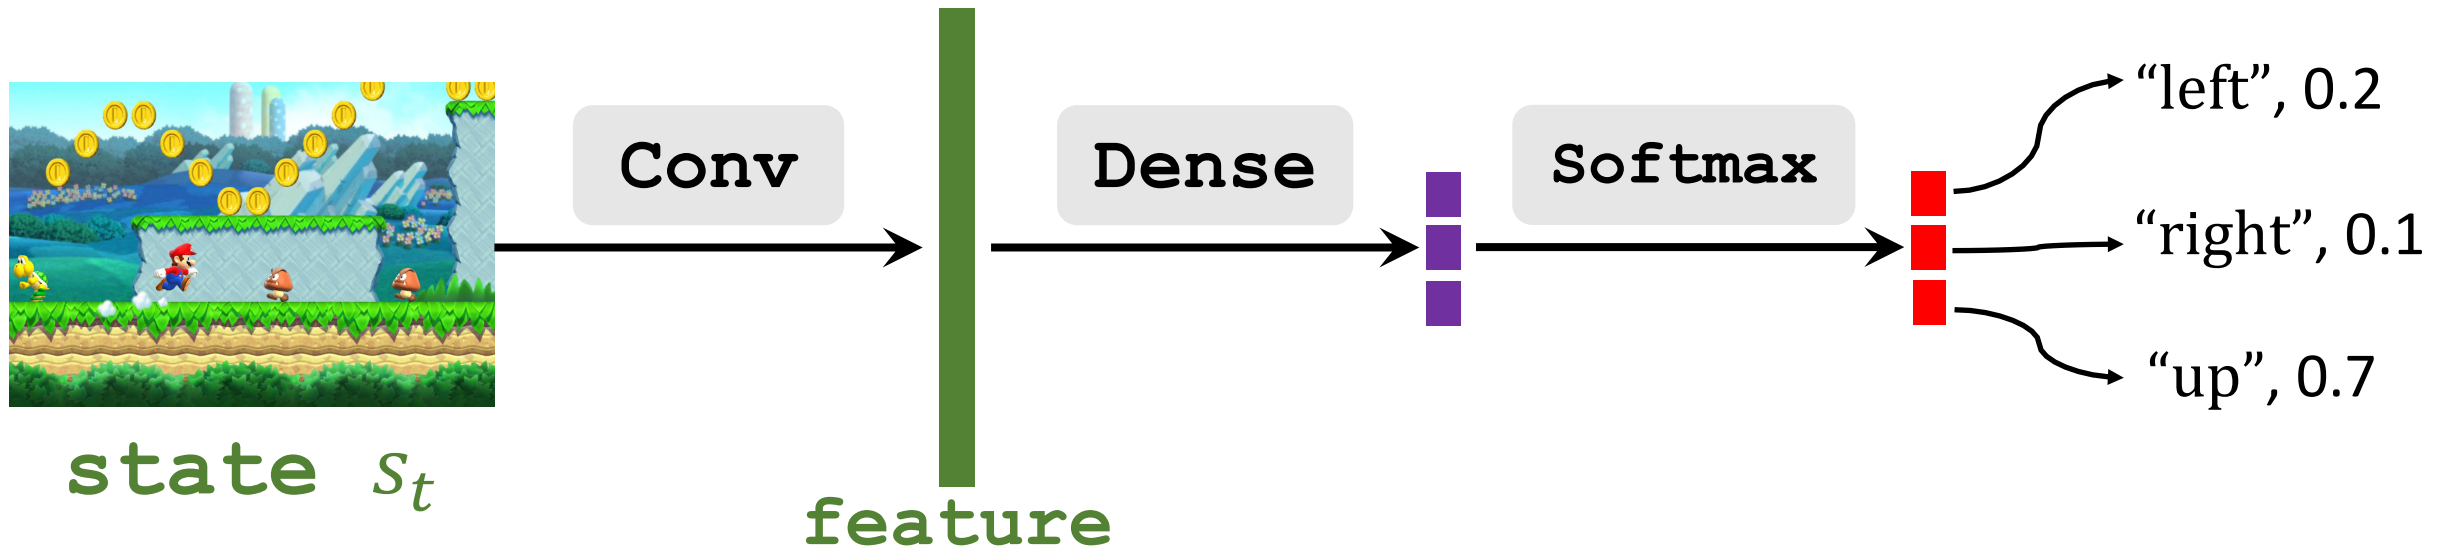
**Policy network:** Use a neural net to approximate  $\pi(a|s)$ .

- Use policy network  $\pi(a|s; \theta)$  to approximate  $\pi(a|s)$ .
- $\theta$ : trainable parameters of the neural net.

# Policy Network $\pi(a|s; \theta)$

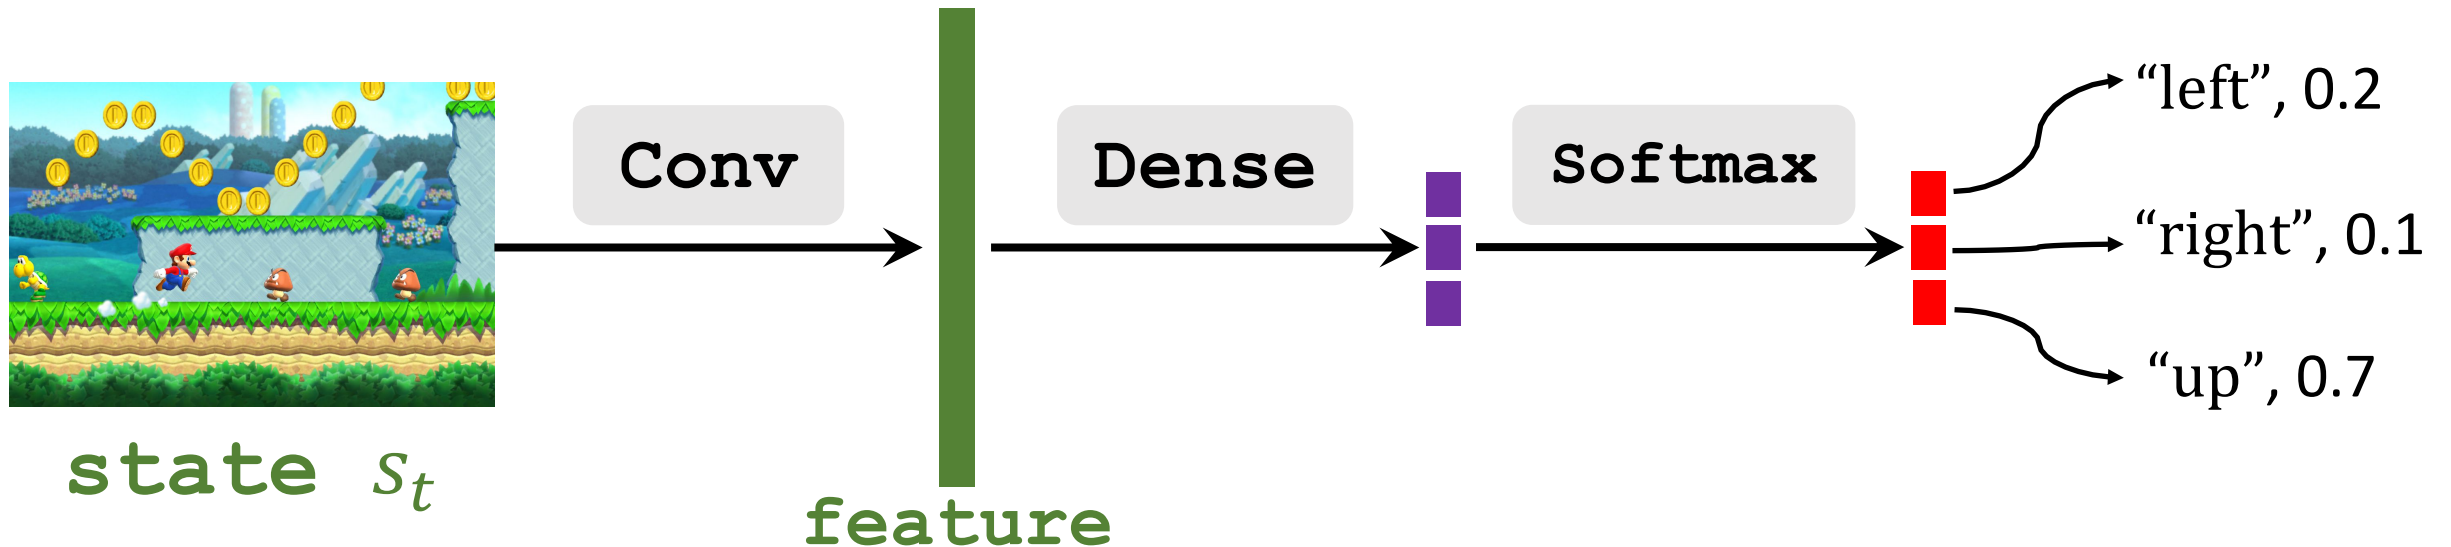
**Policy network:** Use a neural net to approximate  $\pi(a|s)$ .

- Use policy network  $\pi(a|s; \theta)$  to approximate  $\pi(a|s)$ .
- $\theta$ : trainable parameters of the neural net.



# Policy Network $\pi(a|s; \theta)$

- $\sum_{a \in \mathcal{A}} \pi(a|s; \theta) = 1$ .
- Here,  $\mathcal{A} = \{\text{"left"}, \text{"right"}, \text{"up"}\}$  is the set all actions.
- That is why we use softmax activation.





# **State-Value Function Approximation**

# Action-Value Function

**Definition:** Discounted return.

- $U_t = R_t + \gamma \cdot R_{t+1} + \gamma^2 \cdot R_{t+2} + \gamma^3 \cdot R_{t+3} + \dots$



- The return depends on actions  $A_t, A_{t+1}, A_{t+2}, \dots$  and states  $S_t, S_{t+1}, S_{t+2}, \dots$
- Actions are random:  $\mathbb{P}[A = a \mid S = s] = \pi(a|s)$ . (Policy function.)
- States are random:  $\mathbb{P}[S' = s' \mid S = s, A = a] = p(s'|s, a)$ . (State transition.)


# Action-Value Function

**Definition:** Discounted return.

- $U_t = R_t + \gamma \cdot R_{t+1} + \gamma^2 \cdot R_{t+2} + \gamma^3 \cdot R_{t+3} + \dots$

**Definition:** Action-value function.

- $Q_\pi(s_t, a_t) = \mathbb{E} [U_t | S_t = s_t, A_t = a_t].$



The expectation is taken w.r.t.  
actions  $A_{t+1}, A_{t+2}, A_{t+3}, \dots$   
and states  $S_{t+1}, S_{t+2}, S_{t+3}, \dots$

# State-Value Function

**Definition:** Discounted return.

- $U_t = R_t + \gamma \cdot R_{t+1} + \gamma^2 \cdot R_{t+2} + \gamma^3 \cdot R_{t+3} + \dots$

**Definition:** Action-value function.

- $Q_\pi(s_t, a_t) = \mathbb{E} [U_t | S_t = s_t, A_t = a_t].$

**Definition:** State-value function.

- $V_\pi(s_t) = \mathbb{E}_A [Q_\pi(s_t, A)]$



Integrate out action  $A \sim \pi(\cdot | s_t)$ .

# State-Value Function

**Definition:** Discounted return.

- $U_t = R_t + \gamma \cdot R_{t+1} + \gamma^2 \cdot R_{t+2} + \gamma^3 \cdot R_{t+3} + \dots$

**Definition:** Action-value function.

- $Q_\pi(s_t, a_t) = \mathbb{E} [U_t | S_t = s_t, A_t = a_t].$

**Definition:** State-value function.

- $V_\pi(s_t) = \mathbb{E}_A [Q_\pi(s_t, A)] = \sum_a \pi(a | s_t) \cdot Q_\pi(s_t, a).$



Integrate out action  $A \sim \pi(\cdot | s_t).$

# Policy-Based Reinforcement Learning

**Definition:** State-value function.

- $V_{\pi}(s_t) = \mathbb{E}_A [Q_{\pi}(s_t, A)] = \sum_a \pi(a|s_t) \cdot Q_{\pi}(s_t, a).$

# Policy-Based Reinforcement Learning

**Definition:** State-value function.

- $V_{\pi}(s_t) = \mathbb{E}_A [Q_{\pi}(s_t, A)] = \sum_a \pi(a|s_t) \cdot Q_{\pi}(s_t, a).$

Approximate state-value function.

- Approximate policy function  $\pi(a|s_t)$  by policy network  $\pi(a|s_t; \theta).$

# Policy-Based Reinforcement Learning

**Definition:** State-value function.

- $V_{\pi}(s_t) = \mathbb{E}_A [Q_{\pi}(s_t, A)] = \sum_a \pi(a|s_t) \cdot Q_{\pi}(s_t, a).$

Approximate state-value function.

- Approximate policy function  $\pi(a|s_t)$  by policy network  $\pi(a|s_t; \theta).$
- Approximate value function  $V_{\pi}(s_t)$  by:

$$V(s_t; \theta) = \sum_a \pi(a|s_t; \theta) \cdot Q_{\pi}(s_t, a).$$



# Policy-Based Reinforcement Learning

**Definition:** Approximate state-value function.

- $V(s; \theta) = \sum_a \pi(a|s; \theta) \cdot Q_\pi(s, a).$

**Policy-based learning:** Learn  $\theta$  that maximizes  $J(\theta) = \mathbb{E}_s[V(S; \theta)].$

# Policy-Based Reinforcement Learning

**Definition:** Approximate state-value function.

- $V(s; \theta) = \sum_a \pi(a|s; \theta) \cdot Q_\pi(s, a).$

**Policy-based learning:** Learn  $\theta$  that maximizes  $J(\theta) = \mathbb{E}_s[V(S; \theta)].$

How to improve  $\theta$ ? Policy gradient ascent!

- Observe state  $s$ .

- Update policy by:  $\theta \leftarrow \theta + \beta \cdot \frac{\partial V(s; \theta)}{\partial \theta}$

Policy gradient

# Policy Gradient

## Reference

- Sutton and others: [Policy gradient methods for reinforcement learning with function approximation](#). In *NIPS*, 2000.

# Policy Gradient

**Definition:** Approximate state-value function.

- $V(s; \theta) = \sum_a \pi(a|s; \theta) \cdot Q_\pi(s, a).$

**Policy gradient:** Derivative of  $V(s; \theta)$  w.r.t.  $\theta$ .

- $\frac{\partial V(s; \theta)}{\partial \theta}$

# Policy Gradient

**Definition:** Approximate state-value function.

- $V(s; \theta) = \sum_a \pi(a|s; \theta) \cdot Q_\pi(s, a).$

**Policy gradient:** Derivative of  $V(s; \theta)$  w.r.t.  $\theta$ .

- $$\frac{\partial V(s; \theta)}{\partial \theta} = \frac{\partial \sum_a \pi(a|s; \theta) \cdot Q_\pi(s, a)}{\partial \theta}$$

# Policy Gradient

**Definition:** Approximate state-value function.

- $V(s; \theta) = \sum_a \pi(a|s; \theta) \cdot Q_\pi(s, a).$

**Policy gradient:** Derivative of  $V(s; \theta)$  w.r.t.  $\theta$ .

- $$\frac{\partial V(s; \theta)}{\partial \theta} = \frac{\partial \sum_a \pi(a|s; \theta) \cdot Q_\pi(s, a)}{\partial \theta}$$
$$= \sum_a \frac{\partial \pi(a|s; \theta) \cdot Q_\pi(s, a)}{\partial \theta}$$

← Push derivative inside the summation

# Policy Gradient

**Definition:** Approximate state-value function.

- $V(s; \theta) = \sum_a \pi(a|s; \theta) \cdot Q_\pi(s, a).$

**Policy gradient:** Derivative of  $V(s; \theta)$  w.r.t.  $\theta$ .

- $$\begin{aligned} \frac{\partial V(s; \theta)}{\partial \theta} &= \frac{\partial \sum_a \pi(a|s; \theta) \cdot Q_\pi(s, a)}{\partial \theta} \\ &= \sum_a \frac{\partial \pi(a|s; \theta) \cdot Q_\pi(s, a)}{\partial \theta} \\ &= \sum_a \frac{\partial \pi(a|s; \theta)}{\partial \theta} \cdot Q_\pi(s, a) \end{aligned}$$

← Pretend  $Q_\pi$  is independent of  $\theta$ .  
(It may not be true.)

# Policy Gradient

**Definition:** Approximate state-value function.

- $V(s; \theta) = \sum_a \pi(a|s; \theta) \cdot Q_\pi(s, a).$

**Policy gradient:** Derivative of  $V(s; \theta)$  w.r.t.  $\theta$ .

- $\frac{\partial V(s; \theta)}{\partial \theta} = \sum_a \frac{\partial \pi(a|s; \theta)}{\partial \theta} \cdot Q_\pi(s, a)$

Policy Gradient



# Policy Gradient

**Definition:** Approximate state-value function.

- $V(s; \theta) = \sum_a \pi(a|s; \theta) \cdot Q_\pi(s, a).$

**Policy gradient:** Derivative of  $V(s; \theta)$  w.r.t.  $\theta$ .

- $\frac{\partial V(s; \theta)}{\partial \theta} = \sum_a \frac{\partial \pi(a|s; \theta)}{\partial \theta} \cdot Q_\pi(s, a)$

Policy Gradient

**Note:** This derivation is **over-simplified** and **not rigorous**.

# Policy Gradient

**Definition:** Approximate state-value function.

- $V(s; \theta) = \sum_a \pi(a|s; \theta) \cdot Q_\pi(s, a).$

**Policy gradient:** Derivative of  $V(s; \theta)$  w.r.t.  $\theta$ .

- $$\begin{aligned} \frac{\partial V(s; \theta)}{\partial \theta} &= \sum_a \frac{\partial \pi(a|s; \theta)}{\partial \theta} Q_\pi(s, a) \\ &= \sum_a \pi(a|s; \theta) \cdot \frac{\partial \log \pi(a|s; \theta)}{\partial \theta} \cdot Q_\pi(s, a) \end{aligned}$$

# Policy Gradient

**Definition:** Approximate state-value function.

- $V(s; \theta) = \sum_a \pi(a|s; \theta) \cdot Q_\pi(s, a).$

**Policy gradient:** Derivative of  $V(s; \theta)$  w.r.t.  $\theta$ .

- $$\begin{aligned} \frac{\partial V(s; \theta)}{\partial \theta} &= \sum_a \frac{\partial \pi(a|s; \theta)}{\partial \theta} \cdot Q_\pi(s, a) \\ &= \sum_a \pi(a|s; \theta) \cdot \boxed{\frac{\partial \log \pi(a|s; \theta)}{\partial \theta}} \cdot Q_\pi(s, a) \end{aligned}$$

- Chain rule:  $\frac{\partial \log[\pi(\theta)]}{\partial \theta} = \frac{1}{\pi(\theta)} \cdot \frac{\partial \pi(\theta)}{\partial \theta}.$

# Policy Gradient

**Definition:** Approximate state-value function.

- $V(s; \theta) = \sum_a \pi(a|s; \theta) \cdot Q_\pi(s, a).$

**Policy gradient:** Derivative of  $V(s; \theta)$  w.r.t.  $\theta$ .

- $$\begin{aligned} \frac{\partial V(s; \theta)}{\partial \theta} &= \sum_a \frac{\partial \pi(a|s; \theta)}{\partial \theta} \cdot Q_\pi(s, a) \\ &= \sum_a \pi(a|s; \theta) \cdot \frac{\partial \log \pi(a|s; \theta)}{\partial \theta} \cdot Q_\pi(s, a) \end{aligned}$$

- Chain rule:  $\frac{\partial \log[\pi(\theta)]}{\partial \theta} = \frac{1}{\pi(\theta)} \cdot \frac{\partial \pi(\theta)}{\partial \theta}.$
- $\rightarrow \pi(\theta) \cdot \frac{\partial \log[\pi(\theta)]}{\partial \theta} = \pi(\theta) \cdot \frac{1}{\pi(\theta)} \cdot \frac{\partial \pi(\theta)}{\partial \theta}$

# Policy Gradient

**Definition:** Approximate state-value function.

- $V(s; \theta) = \sum_a \pi(a|s; \theta) \cdot Q_\pi(s, a).$

**Policy gradient:** Derivative of  $V(s; \theta)$  w.r.t.  $\theta$ .

- $$\begin{aligned} \frac{\partial V(s; \theta)}{\partial \theta} &= \sum_a \frac{\partial \pi(a|s; \theta)}{\partial \theta} \cdot Q_\pi(s, a) \\ &= \sum_a \pi(a|s; \theta) \cdot \frac{\partial \log \pi(a|s; \theta)}{\partial \theta} \cdot Q_\pi(s, a) \end{aligned}$$

- Chain rule:  $\frac{\partial \log[\pi(\theta)]}{\partial \theta} = \frac{1}{\pi(\theta)} \cdot \frac{\partial \pi(\theta)}{\partial \theta}.$
- $\rightarrow \pi(\theta) \cdot \frac{\partial \log[\pi(\theta)]}{\partial \theta} = \cancel{\pi(\theta)} \cdot \cancel{\frac{1}{\pi(\theta)}} \cdot \frac{\partial \pi(\theta)}{\partial \theta}.$

# Policy Gradient

**Definition:** Approximate state-value function.

- $V(s; \theta) = \sum_a \pi(a|s; \theta) \cdot Q_\pi(s, a).$

**Policy gradient:** Derivative of  $V(s; \theta)$  w.r.t.  $\theta$ .

- $$\begin{aligned} \frac{\partial V(s; \theta)}{\partial \theta} &= \sum_a \frac{\partial \pi(a|s; \theta)}{\partial \theta} Q_\pi(s, a) \\ &= \sum_a \pi(a|s; \theta) \cdot \frac{\partial \log \pi(a|s; \theta)}{\partial \theta} \cdot Q_\pi(s, a) \end{aligned}$$

- Chain rule:  $\frac{\partial \log[\pi(\theta)]}{\partial \theta} = \frac{1}{\pi(\theta)} \cdot \frac{\partial \pi(\theta)}{\partial \theta}.$
- $\rightarrow \pi(\theta) \cdot \frac{\partial \log[\pi(\theta)]}{\partial \theta} = \cancel{\pi(\theta)} \cdot \frac{1}{\cancel{\pi(\theta)}} \cdot \frac{\partial \pi(\theta)}{\partial \theta}.$

# Policy Gradient

**Definition:** Approximate state-value function.

- $V(\mathbf{s}; \boldsymbol{\theta}) = \sum_{\mathbf{a}} \pi(\mathbf{a}|\mathbf{s}; \boldsymbol{\theta}) \cdot Q_{\pi}(\mathbf{s}, \mathbf{a}).$

**Policy gradient:** Derivative of  $V(\mathbf{s}; \boldsymbol{\theta})$  w.r.t.  $\boldsymbol{\theta}$ .

- $$\begin{aligned} \frac{\partial V(\mathbf{s}; \boldsymbol{\theta})}{\partial \boldsymbol{\theta}} &= \sum_{\mathbf{a}} \frac{\partial \pi(\mathbf{a}|\mathbf{s}; \boldsymbol{\theta})}{\partial \boldsymbol{\theta}} \cdot Q_{\pi}(\mathbf{s}, \mathbf{a}) \\ &= \sum_{\mathbf{a}} \pi(\mathbf{a}|\mathbf{s}; \boldsymbol{\theta}) \cdot \frac{\partial \log \pi(\mathbf{a}|\mathbf{s}; \boldsymbol{\theta})}{\partial \boldsymbol{\theta}} \cdot Q_{\pi}(\mathbf{s}, \mathbf{a}) \\ &= \mathbb{E}_{\mathbf{A}} \left[ \frac{\partial \log \pi(\mathbf{A}|\mathbf{s}; \boldsymbol{\theta})}{\partial \boldsymbol{\theta}} \cdot Q_{\pi}(\mathbf{s}, \mathbf{A}) \right]. \end{aligned}$$

The expectation is taken w.r.t. the random variable  $\mathbf{A} \sim \pi(\cdot | \mathbf{s}; \boldsymbol{\theta})$ .

# Policy Gradient

Policy gradient:

$$\frac{\partial V(\mathbf{s}; \boldsymbol{\theta})}{\partial \boldsymbol{\theta}} = \mathbb{E}_{\mathbf{A} \sim \pi(\cdot | \mathbf{s}; \boldsymbol{\theta})} \left[ \frac{\partial \log \pi(\mathbf{A} | \mathbf{s}, \boldsymbol{\theta})}{\partial \boldsymbol{\theta}} \cdot Q_{\pi}(\mathbf{s}, \mathbf{A}) \right].$$



# Calculate Policy Gradient

Policy Gradient: 
$$\frac{\partial V(\mathbf{s}; \boldsymbol{\theta})}{\partial \boldsymbol{\theta}} = \mathbb{E}_{\mathbf{A} \sim \pi(\cdot | \mathbf{s}; \boldsymbol{\theta})} \left[ \frac{\partial \log \pi(\mathbf{A} | \mathbf{s}, \boldsymbol{\theta})}{\partial \boldsymbol{\theta}} \cdot Q_{\pi}(\mathbf{s}, \mathbf{A}) \right].$$

# Calculate Policy Gradient

**Policy Gradient:** 
$$\frac{\partial V(\mathbf{s}; \boldsymbol{\theta})}{\partial \boldsymbol{\theta}} = \mathbb{E}_{\mathbf{A} \sim \pi(\cdot | \mathbf{s}; \boldsymbol{\theta})} \left[ \frac{\partial \log \pi(\mathbf{A} | \mathbf{s}, \boldsymbol{\theta})}{\partial \boldsymbol{\theta}} \cdot Q_{\pi}(\mathbf{s}, \mathbf{A}) \right].$$

1. Randomly sample an action  $\hat{\mathbf{a}}$  according to  $\pi(\cdot | \mathbf{s}; \boldsymbol{\theta})$ .

# Calculate Policy Gradient

**Policy Gradient:**  $\frac{\partial V(s; \theta)}{\partial \theta} = \mathbb{E}_{A \sim \pi(\cdot | s; \theta)} \left[ \frac{\partial \log \pi(A | s, \theta)}{\partial \theta} \cdot Q_{\pi}(s, A) \right].$

1. Randomly sample an action  $\hat{a}$  according to  $\pi(\cdot | s; \theta)$ .

2. Calculate  $\mathbf{g}(\hat{a}, \theta) = \frac{\partial \log \pi(\hat{a} | s; \theta)}{\partial \theta} \cdot Q_{\pi}(s, \hat{a}).$

- By the definition of  $\mathbf{g}$ ,  $\mathbb{E}_A[\mathbf{g}(A, \theta)] = \frac{\partial V(s; \theta)}{\partial \theta}.$
- $\mathbf{g}(\hat{a}, \theta)$  is an unbiased estimate of  $\frac{\partial V(s; \theta)}{\partial \theta}.$

# Calculate Policy Gradient

**Policy Gradient:**  $\frac{\partial V(\mathbf{s}; \boldsymbol{\theta})}{\partial \boldsymbol{\theta}} = \mathbb{E}_{\mathbf{A} \sim \pi(\cdot | \mathbf{s}; \boldsymbol{\theta})} \left[ \frac{\partial \log \pi(\mathbf{A} | \mathbf{s}; \boldsymbol{\theta})}{\partial \boldsymbol{\theta}} \cdot Q_{\pi}(\mathbf{s}, \mathbf{A}) \right].$

1. Randomly sample an action  $\hat{\mathbf{a}}$  according to  $\pi(\cdot | \mathbf{s}; \boldsymbol{\theta})$ .
2. Calculate  $\mathbf{g}(\hat{\mathbf{a}}, \boldsymbol{\theta}) = \frac{\partial \log \pi(\hat{\mathbf{a}} | \mathbf{s}; \boldsymbol{\theta})}{\partial \boldsymbol{\theta}} \cdot Q_{\pi}(\mathbf{s}, \hat{\mathbf{a}})$ .
3. Use  $\mathbf{g}(\hat{\mathbf{a}}, \boldsymbol{\theta})$  as an approximation to the policy gradient  $\frac{\partial V(\mathbf{s}; \boldsymbol{\theta})}{\partial \boldsymbol{\theta}}$ .

Update **policy network** using **policy gradient**

# Algorithm

1. Observe the state  $s_t$ .
2. Randomly sample action  $a_t$  according to  $\pi(\cdot | s_t; \theta_t)$ .

# Algorithm

1. Observe the state  $s_t$ .
2. Randomly sample action  $a_t$  according to  $\pi(\cdot | s_t; \theta_t)$ .
3. Compute  $q_t \approx Q_\pi(s_t, a_t)$  (some estimate).
4. Differentiate policy network:  $\mathbf{d}_{\theta,t} = \frac{\partial \log \pi(a_t | s_t, \theta)}{\partial \theta} \big|_{\theta=\theta_t}$ .

# Algorithm

1. Observe the state  $s_t$ .
2. Randomly sample action  $a_t$  according to  $\pi(\cdot | s_t; \theta_t)$ .
3. Compute  $q_t \approx Q_\pi(s_t, a_t)$  (some estimate).
4. Differentiate policy network:  $\mathbf{d}_{\theta,t} = \frac{\partial \log \pi(a_t | s_t, \theta)}{\partial \theta} \big|_{\theta=\theta_t}$ .
5. (Approximate) policy gradient:  $\mathbf{g}(a_t, \theta_t) = q_t \cdot \mathbf{d}_{\theta,t}$ .
6. Update policy network:  $\theta_{t+1} = \theta_t + \beta \cdot \mathbf{g}(a_t, \theta_t)$ .



# Algorithm

1. Observe the state  $s_t$ .
2. Randomly sample action  $a_t$  according to  $\pi(\cdot | s_t; \theta_t)$ .
3. Compute  $q_t \approx Q_\pi(s_t, a_t)$  (some estimate). **How?**
4. Differentiate policy network:  $\mathbf{d}_{\theta,t} = \frac{\partial \log \pi(a_t | s_t, \theta)}{\partial \theta} \big|_{\theta=\theta_t}$ .
5. (Approximate) policy gradient:  $\mathbf{g}(a_t, \theta_t) = q_t \cdot \mathbf{d}_{\theta,t}$ .
6. Update policy network:  $\theta_{t+1} = \theta_t + \beta \cdot \mathbf{g}(a_t, \theta_t)$ .

# Algorithm

1. Observe the state  $s_t$ .
2. Randomly sample action  $a_t$  according to  $\pi(\cdot | s_t; \theta_t)$ .
3. Compute  $q_t \approx Q_\pi(s_t, a_t)$  (some estimate). **How?**

## Option 1: REINFORCE.

- Play the game to the end and generate the trajectory:

$$s_1, a_1, r_1, s_2, a_2, r_2, \dots, s_T, a_T, r_T.$$

- Compute the discounted return  $u_t = \sum_{k=t}^T \gamma^{k-t} r_k$ , for all  $t$ .
- Since  $Q_\pi(s_t, a_t) = \mathbb{E}[U_t]$ , we can use  $u_t$  to approximate  $Q_\pi(s_t, a_t)$ .
- $\Rightarrow$  Use  $q_t = u_t$ .

# Algorithm

1. Observe the state  $s_t$ .
2. Randomly sample action  $a_t$  according to  $\pi(\cdot | s_t; \theta_t)$ .
3. Compute  $q_t \approx Q_\pi(s_t, a_t)$  (some estimate). **How?**

**Option 2:** Approximate  $Q_\pi$  using a neural network.

- This leads to the actor-critic method.

# Summary

# Policy-Based Learning

- If a good policy function  $\pi$  is known, the agent can be controlled by the policy: randomly sample  $a_t \sim \pi(\cdot | s_t)$ .
- Approximate policy function  $\pi(a|s)$  by **policy network**  $\pi(a|s; \theta)$ .
- Learn the policy network by **policy gradient algorithm**.
- Policy gradient algorithm learn  $\theta$  that maximizes  $\mathbb{E}_s[V(s; \theta)]$ .

**Thank you!**